

Problem Statement

Context

In the realm of modern finance, businesses encounter the perpetual challenge of managing debt obligations effectively to maintain a favorable credit standing and foster sustainable growth. Investors keenly scrutinize companies capable of navigating financial complexities while ensuring stability and profitability. A pivotal instrument in this evaluation process is the balance sheet, which provides a comprehensive overview of a company's assets, liabilities, and shareholder equity, offering insights into its financial health and operational efficiency. In this context, leveraging available financial data, particularly from preceding fiscal periods, becomes imperative for informed decision-making and strategic planning.

Objective

A renowned credit rating organization wants to develop a Financial Health Assessment Tool. With the help of the tool, it endeavors to empower businesses and investors with a robust mechanism for evaluating the financial well-being and creditworthiness of companies. By harnessing machine learning techniques, the organization aims to analyze historical financial statements and extract pertinent insights to facilitate informed decision-making via the tool. Specifically, the organization foresees facilitating the following with the help of the tool:

1. **Debt Management Analysis:** Identify patterns and trends in debt management practices to assess the ability of businesses to fulfill financial obligations promptly and efficiently, and identify potential cases of default.
2. **Credit Risk Evaluation:** Evaluate credit risk exposure by analyzing liquidity ratios, debt-to-equity ratios, and other key financial indicators to ascertain the likelihood of default and inform investment decisions.

As a part of the data science team in the organization, you have been provided with the financial metrics of different companies. The task is to analyze the data provided and develop a predictive model leveraging machine learning techniques to identify whether a given company will default on its debt repayments in the next two quarters. The predictive model will help the organization anticipate potential challenges with the financial performance of the companies and enable proactive risk mitigation strategies.

Data Dictionary

The data consists of financial metrics from the balance sheets of different companies. The detailed data dictionary is available in the data dictionary file (*FRA_DataDictionary.xlsx*).

Please read the instructions carefully before starting the project.

This is a commented Python Notebook file in which all the instructions and tasks to be performed are mentioned.

- Blanks '_____' are provided in the notebook that needs to be filled with an appropriate code to get the correct result. With every '_____' blank, there is a comment that briefly describes what needs to be filled in the blank space.
- Identify the task to be performed correctly, and only then proceed to write the required code.
- Fill the code wherever asked by the commented lines like "# write your code here" or "# complete the code". Running incomplete code may throw error.
- Please run the codes in a sequential manner from the beginning to avoid any unnecessary errors.
- Add the results/observations (wherever mentioned) derived from the analysis in the presentation and submit the same.

Importing necessary libraries

```
import numpy as np
import pandas as pd
import io

import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline

from sklearn.impute import KNNImputer
from sklearn.preprocessing import StandardScaler
from scipy import stats
from statsmodels.stats.outliers_influence import variance_inflation_factor
from sklearn.model_selection import train_test_split, GridSearchCV
# Train test Split and Grid Search
from sklearn.ensemble import RandomForestClassifier

import statsmodels.api as SM
from sklearn import metrics

from sklearn.metrics import (
    confusion_matrix,
    accuracy_score,
    precision_score,
    recall_score,
    f1_score,
```

```

        roc_curve,
        roc_auc_score
    )

import warnings
warnings.filterwarnings('ignore')

from IPython.core.display import display, HTML
display(HTML('<style>.container { width:90% !important; }<\style>'))

<IPython.core.display.HTML object>

```

Loading the Data

```

from google.colab import drive
drive.mount('/content/drive')

Drive already mounted at /content/drive; to attempt to forcibly
remount, call drive.mount("/content/drive", force_remount=True).

df =
pd.read_csv('/content/drive/MyDrive/Finance_and_Risk_Analytics/project
_coded/CompData_1.csv')

```

Data Overview

```

df.head() ## Complete the code to view top 5 rows of the data
{"type": "dataframe", "variable_name": "df"}

df.tail() ## Complete the code to view last 5 rows of the data
{"type": "dataframe"}

df.shape ## Complete the code to view dimensions of the data
(2058, 58)

df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2058 entries, 0 to 2057
Data columns (total 58 columns):
#   Column                                     Non-Null Count
Dtype
---  ---
-----
0    Co_Code                                2058 non-null

```

| | | |
|---------|---|---------------|
| int64 | | |
| 1 | Co_Name | 2058 non-null |
| object | | |
| 2 | _Operating_Expense_Rate | 2058 non-null |
| float64 | | |
| 3 | _Research_and_development_expense_rate | 2058 non-null |
| float64 | | |
| 4 | _Cash_flow_rate | 2058 non-null |
| float64 | | |
| 5 | _Interest_bearing_debt_interest_rate | 2058 non-null |
| float64 | | |
| 6 | _Tax_rate_A | 2058 non-null |
| float64 | | |
| 7 | _Cash_Flow_Per_Share | 1891 non-null |
| float64 | | |
| 8 | _Per_Share_Net_profit_before_tax_Yuan_ | 2058 non-null |
| float64 | | |
| 9 | _Realized_Sales_Gross_Profit_Growth_Rate | 2058 non-null |
| float64 | | |
| 10 | _Operating_Profit_Growth_Rate | 2058 non-null |
| float64 | | |
| 11 | _Continuous_Net_Profit_Growth_Rate | 2058 non-null |
| float64 | | |
| 12 | _Total_Asset_Growth_Rate | 2058 non-null |
| float64 | | |
| 13 | _Net_Value_Growth_Rate | 2058 non-null |
| float64 | | |
| 14 | _Total_Asset_Return_Growth_Rate_Ratio | 2058 non-null |
| float64 | | |
| 15 | _Cash_Reinvestment_perc | 2058 non-null |
| float64 | | |
| 16 | _Current_Ratio | 2058 non-null |
| float64 | | |
| 17 | _Quick_Ratio | 2058 non-null |
| float64 | | |
| 18 | _Interest_Expense_Ratio | 2058 non-null |
| float64 | | |
| 19 | _Total_debt_to_Total_net_worth | 2037 non-null |
| float64 | | |
| 20 | _Long_term_fund_suitability_ratio_A | 2058 non-null |
| float64 | | |
| 21 | _Net_profit_before_tax_to_Paid_in_capital | 2058 non-null |
| float64 | | |
| 22 | _Total_Asset_Turnover | 2058 non-null |
| float64 | | |
| 23 | _Accounts_Receivable_Turnover | 2058 non-null |
| float64 | | |
| 24 | _Average_Collection_Days | 2058 non-null |
| float64 | | |

| | | | |
|---------|--|------|----------|
| 25 | _Inventory_Turnover_Rate_times | 2058 | non-null |
| float64 | | | |
| 26 | _Fixed_Assets_Turnover_Frequency | 2058 | non-null |
| float64 | | | |
| 27 | _Net_Worth_Turnover_Rate_times | 2058 | non-null |
| float64 | | | |
| 28 | _Operating_profit_per_person | 2058 | non-null |
| float64 | | | |
| 29 | _Allocation_rate_per_person | 2058 | non-null |
| float64 | | | |
| 30 | _Quick_Assets_to_Total_Assets | 2058 | non-null |
| float64 | | | |
| 31 | _Cash_to_Total_Assets | 1962 | non-null |
| float64 | | | |
| 32 | _Quick_Assets_to_Current_Liability | 2058 | non-null |
| float64 | | | |
| 33 | _Cash_to_Current_Liability | 2058 | non-null |
| float64 | | | |
| 34 | _Operating_Funds_to_Liability | 2058 | non-null |
| float64 | | | |
| 35 | _Inventory_to_Working_Capital | 2058 | non-null |
| float64 | | | |
| 36 | _Inventory_to_Current_Liability | 2058 | non-null |
| float64 | | | |
| 37 | _Long_term_Liability_to_Current_Assets | 2058 | non-null |
| float64 | | | |
| 38 | _Retained_Earnings_to_Total_Assets | 2058 | non-null |
| float64 | | | |
| 39 | _Total_income_to_Total_expense | 2058 | non-null |
| float64 | | | |
| 40 | _Total_expense_to_Assets | 2058 | non-null |
| float64 | | | |
| 41 | _Current_Asset_Turnover_Rate | 2058 | non-null |
| float64 | | | |
| 42 | _Quick_Asset_Turnover_Rate | 2058 | non-null |
| float64 | | | |
| 43 | _Cash_Turnover_Rate | 2058 | non-null |
| float64 | | | |
| 44 | _Fixed_Assets_to_Assets | 2058 | non-null |
| float64 | | | |
| 45 | _Cash_Flow_to_Total_Assets | 2058 | non-null |
| float64 | | | |
| 46 | _Cash_Flow_to_Liability | 2058 | non-null |
| float64 | | | |
| 47 | _CF0_to_Assets | 2058 | non-null |
| float64 | | | |
| 48 | _Cash_Flow_to_Equity | 2058 | non-null |
| float64 | | | |
| 49 | _Current_Liability_to_Current_Assets | 2044 | non-null |

```

float64
50  _Liability_Assets_Flag                2058 non-null
int64
51  _Total_assets_to_GNP_price            2058 non-null
float64
52  _No_credit_Interval                  2058 non-null
float64
53  _Degree_of_Financial_Leverage_DFL    2058 non-null
float64
54  _Interest_Coverage_Ratio_Interest_expense_to_EBIT 2058 non-null
float64
55  _Net_Income_Flag                    2058 non-null
int64
56  _Equity_to_Liability                 2058 non-null
float64
57  Default                             2058 non-null
int64
dtypes: float64(53), int64(4), object(1)
memory usage: 932.7+ KB

```

```

# Remove '_' (startswith) from column headers where present
for col in df.columns:
    if col.startswith('_'):
        df.rename(columns={col: col[1:]}, inplace=True)

```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 2058 entries, 0 to 2057
```

```
Data columns (total 58 columns):
```

| # | Column | Non-Null Count |
|---|---------------------------------------|----------------|
| 0 | Co_Code | 2058 non-null |
| 1 | Co_Name | 2058 non-null |
| 2 | Operating_Expense_Rate | 2058 non-null |
| 3 | Research_and_development_expense_rate | 2058 non-null |
| 4 | Cash_flow_rate | 2058 non-null |
| 5 | Interest_bearing_debt_interest_rate | 2058 non-null |
| 6 | Tax_rate_A | 2058 non-null |
| 7 | Cash_Flow_Per_Share | 1891 non-null |

| | | | |
|---------|--|------|----------|
| 8 | Per_Share_Net_profit_before_tax_Yuan_ | 2058 | non-null |
| float64 | | | |
| 9 | Realized_Sales_Gross_Profit_Growth_Rate | 2058 | non-null |
| float64 | | | |
| 10 | Operating_Profit_Growth_Rate | 2058 | non-null |
| float64 | | | |
| 11 | Continuous_Net_Profit_Growth_Rate | 2058 | non-null |
| float64 | | | |
| 12 | Total_Asset_Growth_Rate | 2058 | non-null |
| float64 | | | |
| 13 | Net_Value_Growth_Rate | 2058 | non-null |
| float64 | | | |
| 14 | Total_Asset_Return_Growth_Rate_Ratio | 2058 | non-null |
| float64 | | | |
| 15 | Cash_Reinvestment_perc | 2058 | non-null |
| float64 | | | |
| 16 | Current_Ratio | 2058 | non-null |
| float64 | | | |
| 17 | Quick_Ratio | 2058 | non-null |
| float64 | | | |
| 18 | Interest_Expense_Ratio | 2058 | non-null |
| float64 | | | |
| 19 | Total_debt_to_Total_net_worth | 2037 | non-null |
| float64 | | | |
| 20 | Long_term_fund_suitability_ratio_A | 2058 | non-null |
| float64 | | | |
| 21 | Net_profit_before_tax_to_Paid_in_capital | 2058 | non-null |
| float64 | | | |
| 22 | Total_Asset_Turnover | 2058 | non-null |
| float64 | | | |
| 23 | Accounts_Receivable_Turnover | 2058 | non-null |
| float64 | | | |
| 24 | Average_Collection_Days | 2058 | non-null |
| float64 | | | |
| 25 | Inventory_Turnover_Rate_times | 2058 | non-null |
| float64 | | | |
| 26 | Fixed_Assets_Turnover_Frequency | 2058 | non-null |
| float64 | | | |
| 27 | Net_Worth_Turnover_Rate_times | 2058 | non-null |
| float64 | | | |
| 28 | Operating_profit_per_person | 2058 | non-null |
| float64 | | | |
| 29 | Allocation_rate_per_person | 2058 | non-null |
| float64 | | | |
| 30 | Quick_Assets_to_Total_Assets | 2058 | non-null |
| float64 | | | |
| 31 | Cash_to_Total_Assets | 1962 | non-null |
| float64 | | | |
| 32 | Quick_Assets_to_Current_Liability | 2058 | non-null |

| | | |
|---------|--|---------------|
| float64 | | |
| 33 | Cash_to_Current_Liability | 2058 non-null |
| float64 | | |
| 34 | Operating_Funds_to_Liability | 2058 non-null |
| float64 | | |
| 35 | Inventory_to_Working_Capital | 2058 non-null |
| float64 | | |
| 36 | Inventory_to_Current_Liability | 2058 non-null |
| float64 | | |
| 37 | Long_term_Liability_to_Current_Assets | 2058 non-null |
| float64 | | |
| 38 | Retained_Earnings_to_Total_Assets | 2058 non-null |
| float64 | | |
| 39 | Total_income_to_Total_expense | 2058 non-null |
| float64 | | |
| 40 | Total_expense_to_Assets | 2058 non-null |
| float64 | | |
| 41 | Current_Asset_Turnover_Rate | 2058 non-null |
| float64 | | |
| 42 | Quick_Asset_Turnover_Rate | 2058 non-null |
| float64 | | |
| 43 | Cash_Turnover_Rate | 2058 non-null |
| float64 | | |
| 44 | Fixed_Assets_to_Assets | 2058 non-null |
| float64 | | |
| 45 | Cash_Flow_to_Total_Assets | 2058 non-null |
| float64 | | |
| 46 | Cash_Flow_to_Liability | 2058 non-null |
| float64 | | |
| 47 | CF0_to_Assets | 2058 non-null |
| float64 | | |
| 48 | Cash_Flow_to_Equity | 2058 non-null |
| float64 | | |
| 49 | Current_Liability_to_Current_Assets | 2044 non-null |
| float64 | | |
| 50 | Liability_Assets_Flag | 2058 non-null |
| int64 | | |
| 51 | Total_assets_to_GNP_price | 2058 non-null |
| float64 | | |
| 52 | No_credit_Interval | 2058 non-null |
| float64 | | |
| 53 | Degree_of_Financial_Leverage_DFL | 2058 non-null |
| float64 | | |
| 54 | Interest_Coverage_Ratio_Interest_expense_to_EBIT | 2058 non-null |
| float64 | | |
| 55 | Net_Income_Flag | 2058 non-null |
| int64 | | |
| 56 | Equity_to_Liability | 2058 non-null |
| float64 | | |


```

57 Default 2058 non-null
int64
dtypes: float64(53), int64(4), object(1)
memory usage: 932.7+ KB

# checking for duplicate values
df.duplicated().sum() ## Complete the code to check duplicate entries
in the data

np.int64(0)

df.nunique() ## complete the code to check unique entries in the data

```

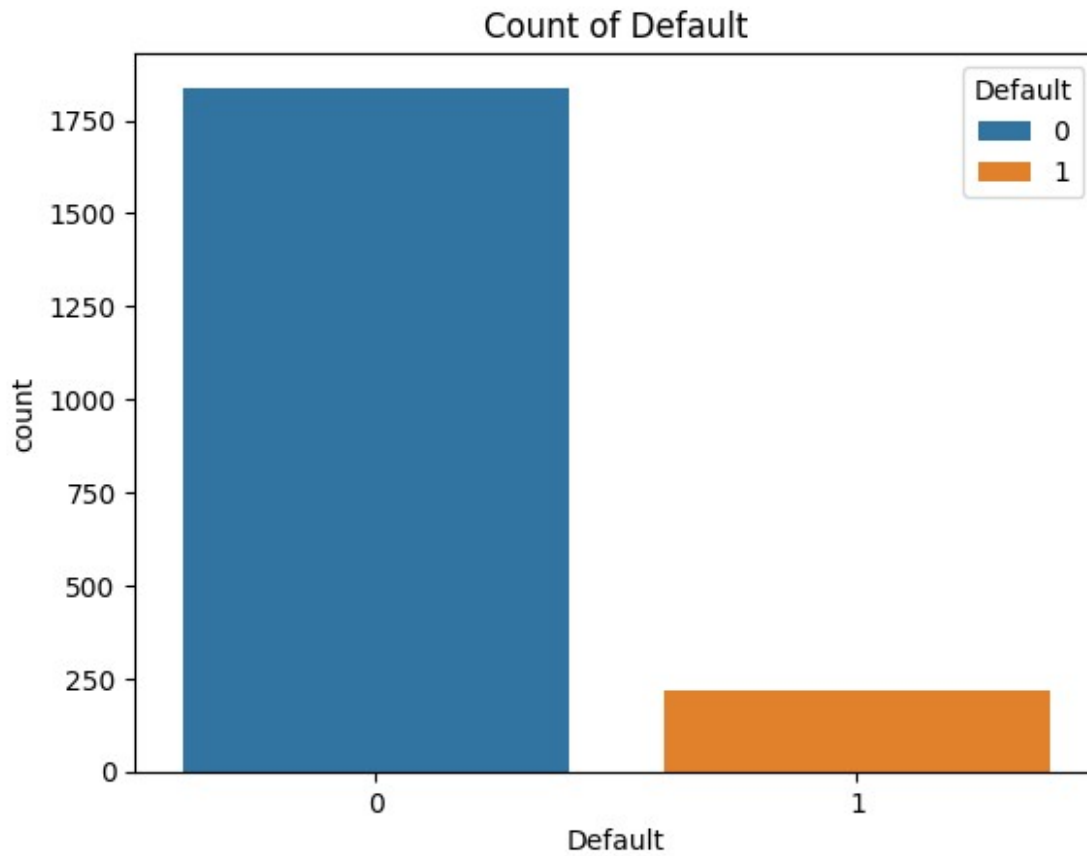
| | |
|--|------|
| Co_Code | 2058 |
| Co_Name | 2058 |
| Operating_Expense_Rate | 1495 |
| Research_and_development_expense_rate | 629 |
| Cash_flow_rate | 1888 |
| Interest_bearing_debt_interest_rate | 813 |
| Tax_rate_A | 985 |
| Cash_Flow_Per_Share | 900 |
| Per_Share_Net_profit_before_tax_Yuan_ | 876 |
| Realized_Sales_Gross_Profit_Growth_Rate | 1939 |
| Operating_Profit_Growth_Rate | 2015 |
| Continuous_Net_Profit_Growth_Rate | 2014 |
| Total_Asset_Growth_Rate | 922 |
| Net_Value_Growth_Rate | 1757 |
| Total_Asset_Return_Growth_Rate_Ratio | 1428 |
| Cash_Reinvestment_perc | 1690 |
| Current_Ratio | 1972 |
| Quick_Ratio | 1970 |
| Interest_Expense_Ratio | 1716 |
| Total_debt_to_Total_net_worth | 1949 |
| Long_term_fund_suitability_ratio_A | 2014 |
| Net_profit_before_tax_to_Paid_in_capital | 1798 |
| Total_Asset_Turnover | 283 |
| Accounts_Receivable_Turnover | 1109 |
| Average_Collection_Days | 1935 |
| Inventory_Turnover_Rate_times | 1151 |
| Fixed_Assets_Turnover_Frequency | 1079 |
| Net_Worth_Turnover_Rate_times | 529 |
| Operating_profit_per_person | 1484 |
| Allocation_rate_per_person | 2051 |
| Quick_Assets_to_Total_Assets | 2058 |
| Cash_to_Total_Assets | 1962 |
| Quick_Assets_to_Current_Liability | 2058 |
| Cash_to_Current_Liability | 2056 |
| Operating_Funds_to_Liability | 2058 |
| Inventory_to_Working_Capital | 1931 |
| Inventory_to_Current_Liability | 1932 |

| | |
|--|------|
| Long_term_Liability_to_Current_Assets | 1398 |
| Retained_Earnings_to_Total_Assets | 2058 |
| Total_income_to_Total_expense | 2056 |
| Total_expense_to_Assets | 2058 |
| Current_Asset_Turnover_Rate | 1973 |
| Quick_Asset_Turnover_Rate | 1743 |
| Cash_Turnover_Rate | 1440 |
| Fixed_Assets_to_Assets | 2054 |
| Cash_Flow_to_Total_Assets | 2058 |
| Cash_Flow_to_Liability | 2058 |
| CF0_to_Assets | 2058 |
| Cash_Flow_to_Equity | 2058 |
| Current_Liability_to_Current_Assets | 2044 |
| Liability_Assets_Flag | 2 |
| Total_assets_to_GNP_price | 2058 |
| No_credit_Interval | 2057 |
| Degree_of_Financial_Leverage_DFL | 1940 |
| Interest_Coverage_Ratio_Interest_expense_to_EBIT | 1945 |
| Net_Income_Flag | 1 |
| Equity_to_Liability | 2058 |
| Default | 2 |

dtype: int64

```
df.describe().T
```

```
{
  "summary": {
    "name": "df",
    "rows": 57,
    "fields": [
      {
        "column": "count",
        "properties": {
          "dtype": "number",
          "std": 25.419352543077746,
          "min": 1891.0,
          "max": 2058.0,
          "num_unique_values": 5,
          "samples": [
            1891.0,
            2044.0,
            2037.0
          ],
          "semantic_type": ""
        },
        "description": ""
      },
      {
        "column": "mean",
        "properties": {
          "dtype": "number",
          "std": 929954049.0062642,
          "min": 0.0023579771039844515,
          "max": 5287663257.045698,
          "num_unique_values": 57,
          "samples": [
            17572.113216715257,
            0.11477699404324587,
            0.07993674991437308
          ],
          "semantic_type": ""
        },
        "description": ""
      },
      {
        "column": "std",
        "properties": {
          "dtype": "number",
          "std": 1005020809.4437544,
          "min": 0.0,
          "max": 3453544121.673858,
          "num_unique_values": 57,
          "samples": [
            21892.886518349056,
            0.15244565360109494,
            0.09862259512732567
          ],
          "semantic_type": ""
        },
        "description": ""
      },
      {
        "column": "min",
        "properties": {
          "dtype": "number",
          "std": 0.5534476428333828,
          "min": 0.0,
          "max": 4.0,
          "num_unique_values": 22,
          "samples": [
            4.0,
            0.026274053,
            0.525126368
          ],
          "semantic_type": ""
        },
        "description": ""
      }
    ]
  }
}
```

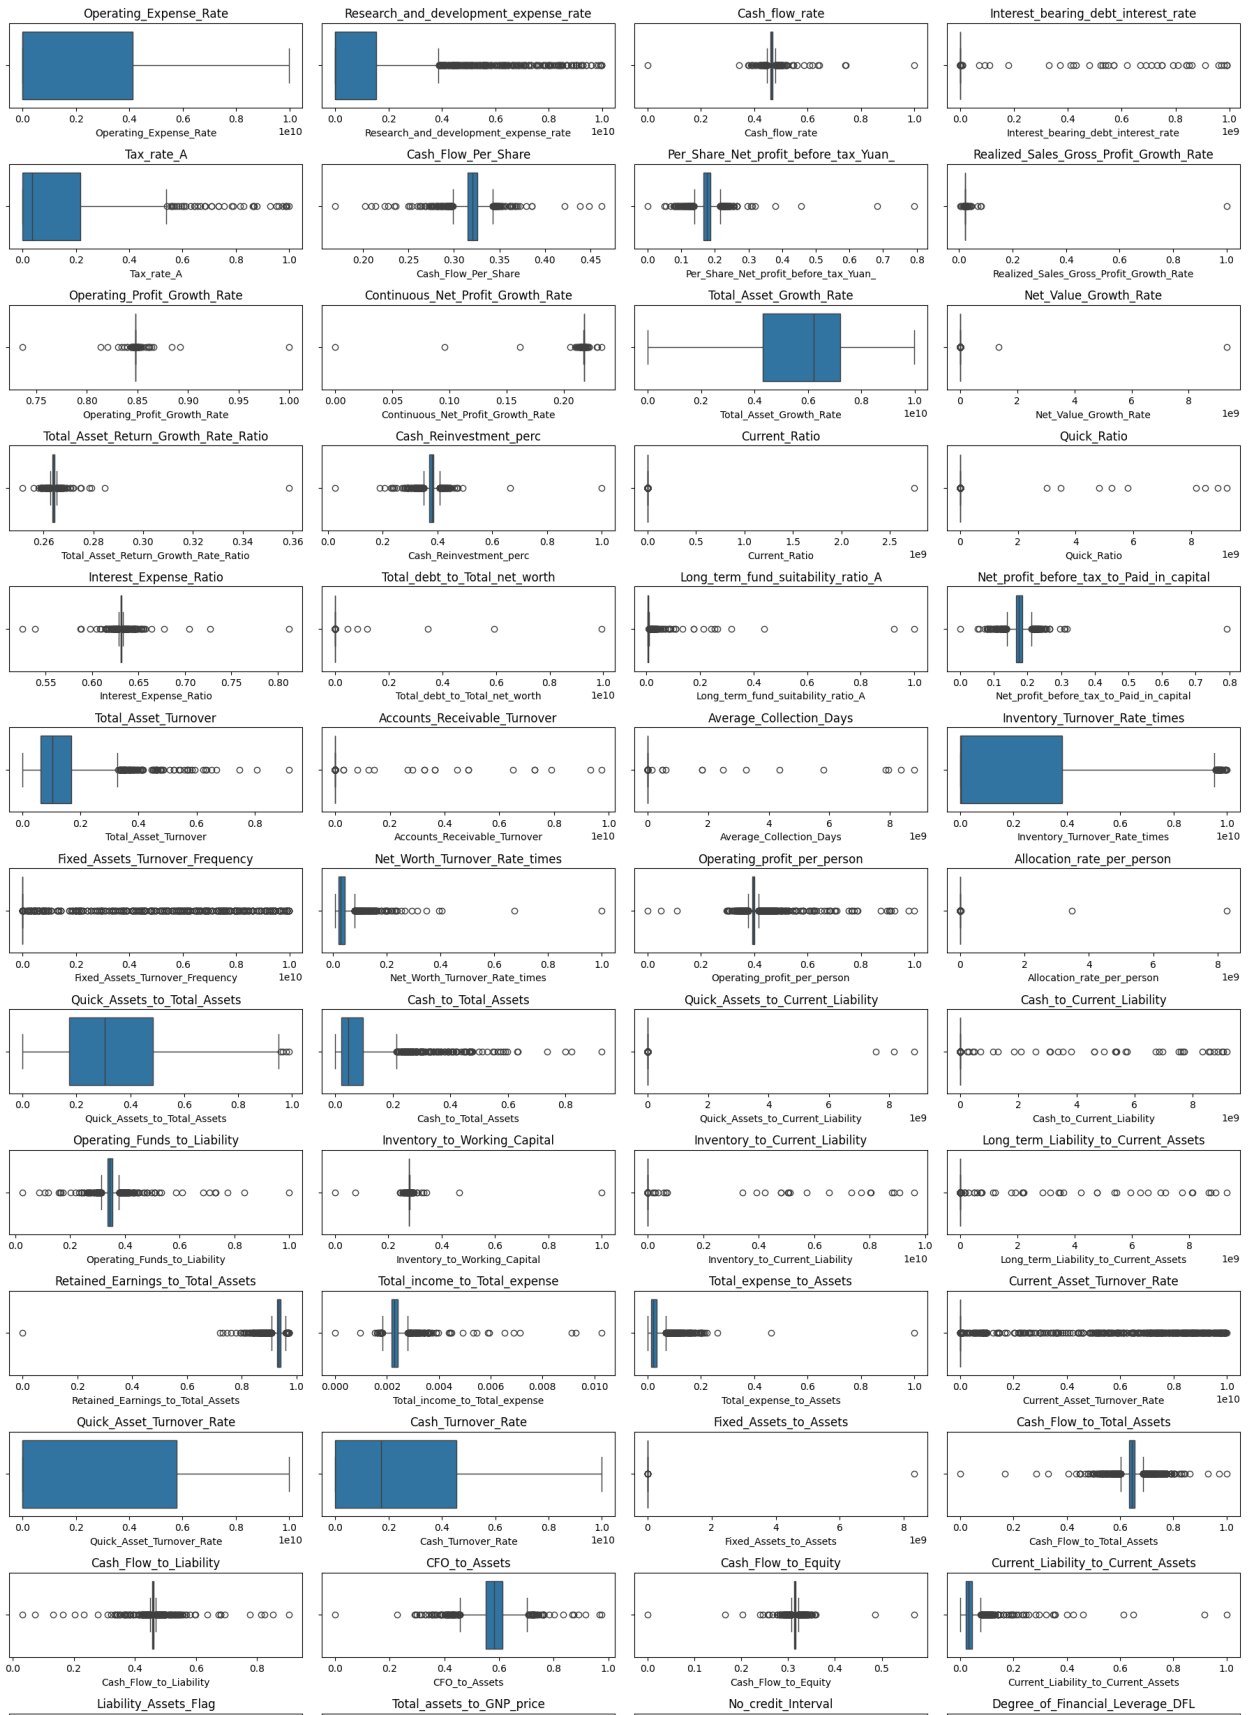
```
#Percentage of defaulters
(df.Default.sum()/len(df)) * 100

np.float64(10.689990281827017)

#Get boxplots for all the numerical columns
numeric_columns = df.select_dtypes(include=np.number).columns.tolist()

plt.figure(figsize=(18, 30))

for i, variable in enumerate(numeric_columns):
    plt.subplot(15, 4, i + 1)
    sns.boxplot(data=df, x=variable) ## Complete the code to get
boxplots for all numerical columns
    plt.tight_layout()
    plt.title(variable)
```



```

#Get distplot for all the numerical columns
numeric_columns = df.select_dtypes(include=np.number).columns.tolist()

plt.figure(figsize=(18, 30))

for i, variable in enumerate(numeric_columns):
    plt.subplot(15, 4, i + 1)
    sns.histplot(data=df, x=variable)  ## Complete the code to get
    histplot for all numerical columns in the data
    plt.tight_layout()
    plt.title(variable)

#Get distplot for all the numerical columns
numeric_columns = df.select_dtypes(include=np.number).columns.tolist()

plt.figure(figsize=(18, 30))

for i, variable in enumerate(numeric_columns):
    plt.subplot(15, 4, i + 1)
    sns.histplot(data=df, x=variable)  ## Complete the code to get
    histplot for all numerical columns in the data
    plt.tight_layout()
    plt.title(variable)

```

Bivariate Analysis

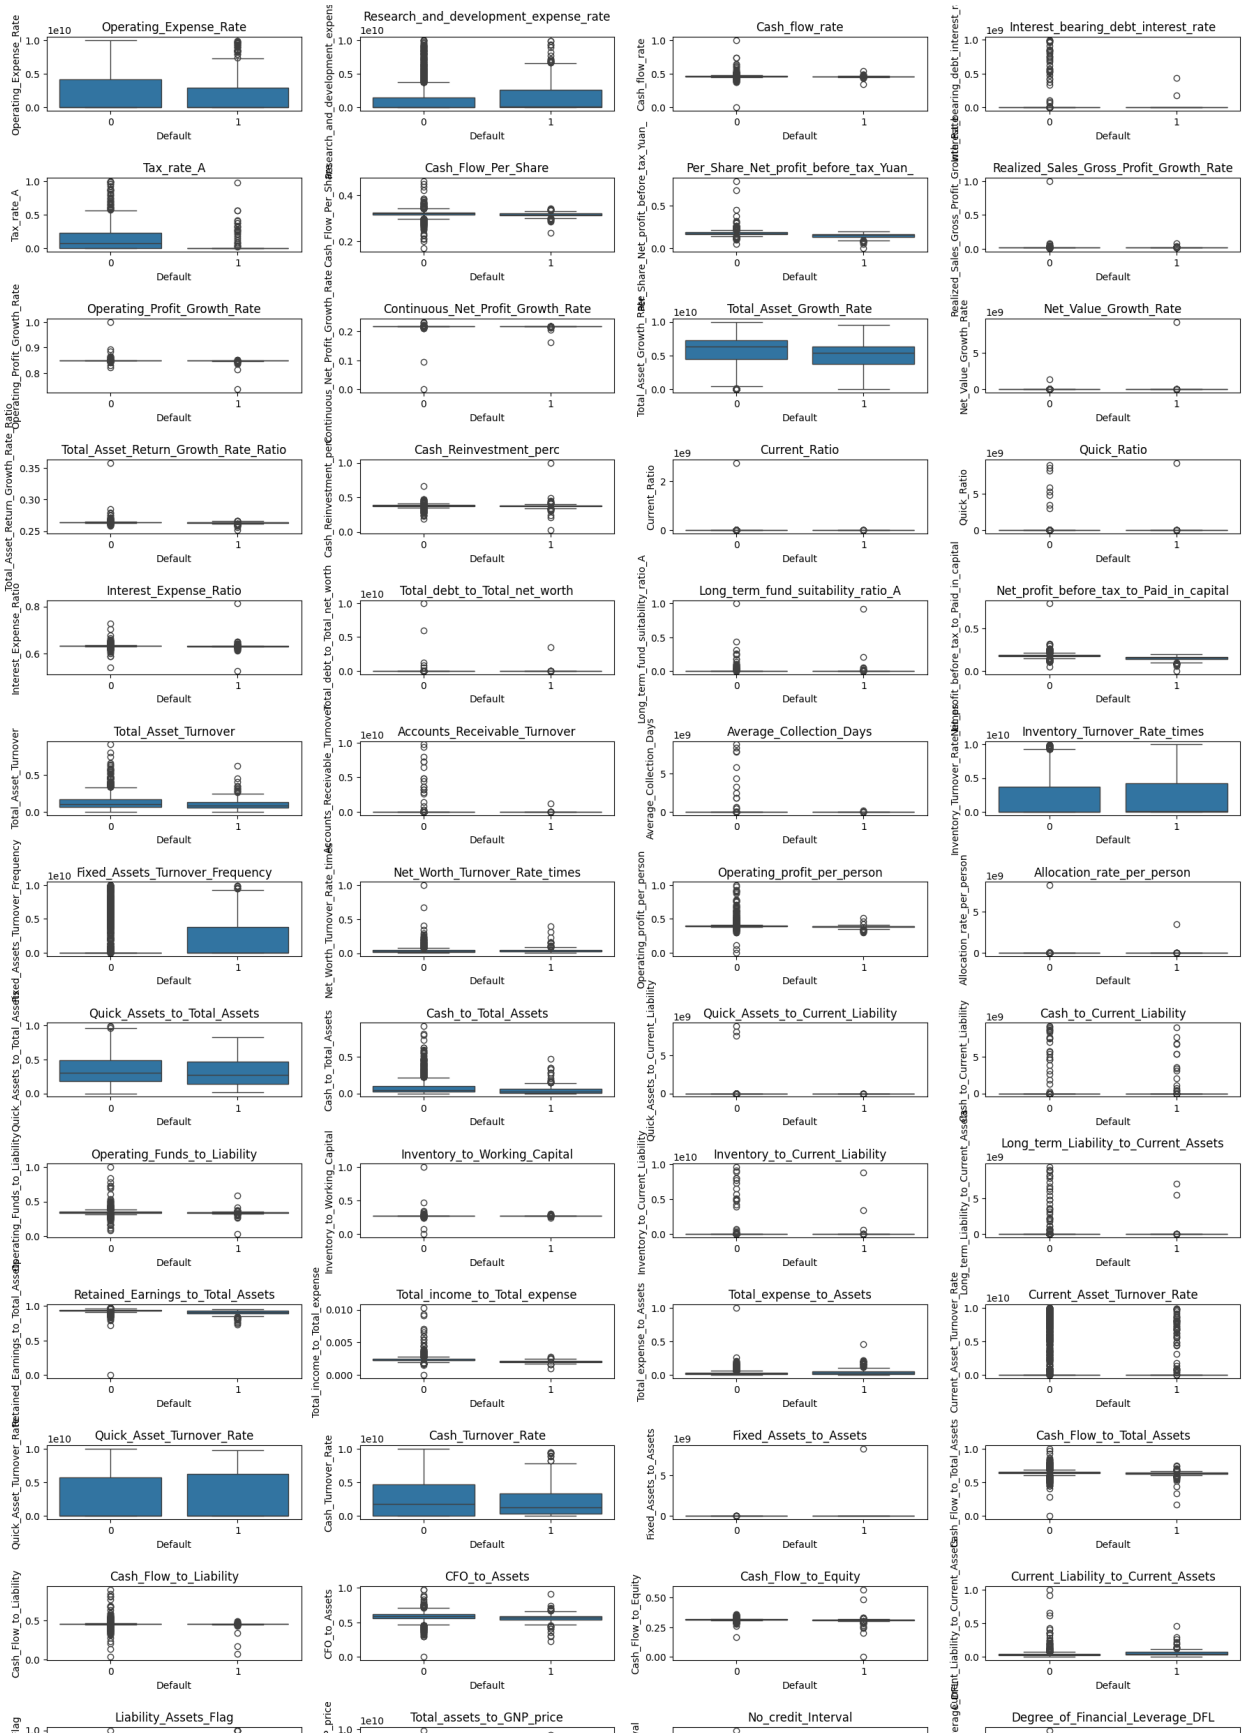
```

#Get boxplots for all the numerical columns
numeric_columns = df.select_dtypes(include=np.number).columns.tolist()

plt.figure(figsize=(18, 30))


for i, variable in enumerate(numeric_columns):
    plt.subplot(15, 4, i + 1)
    sns.boxplot(x="Default", y=variable, data=df)  ## Complete the
    code to get boxplot of all variables with Default column in the data
    plt.tight_layout()
    plt.title(variable)

```



```
# Calculate the correlation matrix
corr_matrix = df.corr(numeric_only=True) ## Complete the code to get
the correlation matrix for the data

# Create a heatmap of the correlation matrix
plt.figure(figsize=(50, 50))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', fmt=".1f",
annot_kws={"size": 15})
plt.title('Heatmap of Correlation Matrix')
plt.show()
```

```
df.nunique()
```

| | |
|--|------|
| Operating_Expense_Rate | 1495 |
| Research_and_development_expense_rate | 629 |
| Cash_flow_rate | 1888 |
| Interest_bearing_debt_interest_rate | 813 |
| Tax_rate_A | 985 |
| Cash_Flow_Per_Share | 900 |
| Per_Share_Net_profit_before_tax_Yuan_ | 876 |
| Realized_Sales_Gross_Profit_Growth_Rate | 1939 |
| Operating_Profit_Growth_Rate | 2015 |
| Continuous_Net_Profit_Growth_Rate | 2014 |
| Total_Asset_Growth_Rate | 922 |
| Net_Value_Growth_Rate | 1757 |
| Total_Asset_Return_Growth_Rate_Ratio | 1428 |
| Cash_Reinvestment_perc | 1690 |
| Current_Ratio | 1972 |
| Quick_Ratio | 1970 |
| Interest_Expense_Ratio | 1716 |
| Total_debt_to_Total_net_worth | 1949 |
| Long_term_fund_suitability_ratio_A | 2014 |
| Net_profit_before_tax_to_Paid_in_capital | 1798 |
| Total_Asset_Turnover | 283 |
| Accounts_Receivable_Turnover | 1109 |
| Average_Collection_Days | 1935 |
| Inventory_Turnover_Rate_times | 1151 |
| Fixed_Assets_Turnover_Frequency | 1079 |
| Net_Worth_Turnover_Rate_times | 529 |
| Operating_profit_per_person | 1484 |
| Allocation_rate_per_person | 2051 |
| Quick_Assets_to_Total_Assets | 2058 |
| Cash_to_Total_Assets | 1962 |
| Quick_Assets_to_Current_Liability | 2058 |
| Cash_to_Current_Liability | 2056 |
| Operating_Funds_to_Liability | 2058 |
| Inventory_to_Working_Capital | 1931 |
| Inventory_to_Current_Liability | 1932 |
| Long_term_Liability_to_Current_Assets | 1398 |
| Retained_Earnings_to_Total_Assets | 2058 |
| Total_income_to_Total_expense | 2056 |
| Total_expense_to_Assets | 2058 |
| Current_Asset_Turnover_Rate | 1973 |
| Quick_Asset_Turnover_Rate | 1743 |
| Cash_Turnover_Rate | 1440 |
| Fixed_Assets_to_Assets | 2054 |
| Cash_Flow_to_Total_Assets | 2058 |
| Cash_Flow_to_Liability | 2058 |
| CF0_to_Assets | 2058 |
| Cash_Flow_to_Equity | 2058 |
| Current_Liability_to_Current_Assets | 2044 |
| Liability_Assets_Flag | 2 |
| Total_assets_to_GNP_price | 2058 |

| | |
|--|------|
| No_credit_Interval | 2057 |
| Degree_of_Financial_Leverage_DFL | 1940 |
| Interest_Coverage_Ratio_Interest_expense_to_EBIT | 1945 |
| Net_Income_Flag | 1 |
| Equity_to_Liability | 2058 |
| Default | 2 |
| dtype: int64 | |

We can drop the columns `Net_Income_Flag` and `Liability_Assets_Flag` as they have very few unique values.

```
df.drop(['Net_Income_Flag', 'Liability_Assets_Flag'], axis = 1,
inplace = True) ## Complete the code to drop the mentioned columns
from the dataset
df.nunique()
```

| | |
|--|------|
| Operating_Expense_Rate | 1495 |
| Research_and_development_expense_rate | 629 |
| Cash_flow_rate | 1888 |
| Interest_bearing_debt_interest_rate | 813 |
| Tax_rate_A | 985 |
| Cash_Flow_Per_Share | 900 |
| Per_Share_Net_profit_before_tax_Yuan | 876 |
| Realized_Sales_Gross_Profit_Growth_Rate | 1939 |
| Operating_Profit_Growth_Rate | 2015 |
| Continuous_Net_Profit_Growth_Rate | 2014 |
| Total_Asset_Growth_Rate | 922 |
| Net_Value_Growth_Rate | 1757 |
| Total_Asset_Return_Growth_Rate_Ratio | 1428 |
| Cash_Reinvestment_perc | 1690 |
| Current_Ratio | 1972 |
| Quick_Ratio | 1970 |
| Interest_Expense_Ratio | 1716 |
| Total_debt_to_Total_net_worth | 1949 |
| Long_term_fund_suitability_ratio_A | 2014 |
| Net_profit_before_tax_to_Paid_in_capital | 1798 |
| Total_Asset_Turnover | 283 |
| Accounts_Receivable_Turnover | 1109 |
| Average_Collection_Days | 1935 |
| Inventory_Turnover_Rate_times | 1151 |
| Fixed_Assets_Turnover_Frequency | 1079 |
| Net_Worth_Turnover_Rate_times | 529 |
| Operating_profit_per_person | 1484 |
| Allocation_rate_per_person | 2051 |
| Quick_Assets_to_Total_Assets | 2058 |
| Cash_to_Total_Assets | 1962 |
| Quick_Assets_to_Current_Liability | 2058 |
| Cash_to_Current_Liability | 2056 |
| Operating_Funds_to_Liability | 2058 |

| | |
|--|------|
| Inventory_to_Working_Capital | 1931 |
| Inventory_to_Current_Liability | 1932 |
| Long_term_Liability_to_Current_Assets | 1398 |
| Retained_Earnings_to_Total_Assets | 2058 |
| Total_income_to_Total_expense | 2056 |
| Total_expense_to_Assets | 2058 |
| Current_Asset_Turnover_Rate | 1973 |
| Quick_Asset_Turnover_Rate | 1743 |
| Cash_Turnover_Rate | 1440 |
| Fixed_Assets_to_Assets | 2054 |
| Cash_Flow_to_Total_Assets | 2058 |
| Cash_Flow_to_Liability | 2058 |
| CF0_to_Assets | 2058 |
| Cash_Flow_to_Equity | 2058 |
| Current_Liability_to_Current_Assets | 2044 |
| Total_assets_to_GNP_price | 2058 |
| No_credit_Interval | 2057 |
| Degree_of_Financial_Leverage_DFL | 1940 |
| Interest_Coverage_Ratio_Interest_expense_to_EBIT | 1945 |
| Equity_to_Liability | 2058 |
| Default | 2 |

dtype: int64

Outliers Check

```

outliers_count = {}

# Iterate over each column in the DataFrame
for column in df.columns:
    # Check if the column is numeric before processing
    if pd.api.types.is_numeric_dtype(df[column]):
        Q1 = df[column].quantile(0.25)
        Q3 = df[column].quantile(0.75)

        IQR = Q3 - Q1

        lower_bound = Q1 - 1.5 * IQR  ## Fill the blank with correct
value for getting lower_bound
        upper_bound = Q3 + 1.5 * IQR  ## Fill the blank with correct
value for getting upper_bound

        outliers = df[(df[column] < lower_bound) | (df[column] >
upper_bound)]
        outliers_count[column] = len(outliers)

print("Number of outliers in each column:")
pd.DataFrame([{'Column': column, 'No. of outliers': outliers} for
column, outliers in outliers_count.items()])

Number of outliers in each column:

```

```
{
  "summary": {
    "name": "pd",
    "rows": 54,
    "fields": [
      {
        "column": "Column",
        "properties": {
          "dtype": "string",
          "num_unique_values": 54,
          "samples": [
            "Net_profit_before_tax_to_Paid_in_capital",
            "No_credit_Interval",
            "Total_assets_to_GNP_price",
            {
              "column": "No. of outliers",
              "properties": {
                "dtype": "number",
                "std": 125,
                "min": 0,
                "max": 501,
                "num_unique_values": 48,
                "samples": [
                  4,
                  407,
                  200
                ],
                "semantic_type": "",
                "description": ""
              }
            }
          ]
        }
      }
    ]
  },
  "type": "dataframe"
}
```

Data Preparation for Modeling

```
# Separating target variable from the rest of the data
df_X = df.drop(['Default'], axis = 1)
df_y = df['Default']

# Splitting the data for training and testing
X_train, X_test, y_train, y_test = train_test_split(df_X, df_y,
  test_size=0.25, random_state=42, stratify = df_y)  ## Complete the
  code to split the data into train and test in the ratio 75:25
```

Missing Values Detection and Treatment

```
# Check missing values
X_train.isnull().sum()  ## Complete the code to get the number of null
  or NaN values in each column
```

| | |
|---|-----|
| Operating_Expense_Rate | 0 |
| Research_and_development_expense_rate | 0 |
| Cash_flow_rate | 0 |
| Interest_bearing_debt_interest_rate | 0 |
| Tax_rate_A | 0 |
| Cash_Flow_Per_Share | 126 |
| Per_Share_Net_profit_before_tax_Yuan_ | 0 |
| Realized_Sales_Gross_Profit_Growth_Rate | 0 |
| Operating_Profit_Growth_Rate | 0 |
| Continuous_Net_Profit_Growth_Rate | 0 |
| Total_Asset_Growth_Rate | 0 |
| Net_Value_Growth_Rate | 0 |
| Total_Asset_Return_Growth_Rate_Ratio | 0 |
| Cash_Reinvestment_perc | 0 |
| Current_Ratio | 0 |
| Quick_Ratio | 0 |
| Interest_Expense_Ratio | 0 |
| Total_debt_to_Total_net_worth | 18 |

| | |
|--|----|
| Long_term_fund_suitability_ratio_A | 0 |
| Net_profit_before_tax_to_Paid_in_capital | 0 |
| Total_Asset_Turnover | 0 |
| Accounts_Receivable_Turnover | 0 |
| Average_Collection_Days | 0 |
| Inventory_Turnover_Rate_times | 0 |
| Fixed_Assets_Turnover_Frequency | 0 |
| Net_Worth_Turnover_Rate_times | 0 |
| Operating_profit_per_person | 0 |
| Allocation_rate_per_person | 0 |
| Quick_Assets_to_Total_Assets | 0 |
| Cash_to_Total_Assets | 71 |
| Quick_Assets_to_Current_Liability | 0 |
| Cash_to_Current_Liability | 0 |
| Operating_Funds_to_Liability | 0 |
| Inventory_to_Working_Capital | 0 |
| Inventory_to_Current_Liability | 0 |
| Long_term_Liability_to_Current_Assets | 0 |
| Retained_Earnings_to_Total_Assets | 0 |
| Total_income_to_Total_expense | 0 |
| Total_expense_to_Assets | 0 |
| Current_Asset_Turnover_Rate | 0 |
| Quick_Asset_Turnover_Rate | 0 |
| Cash_Turnover_Rate | 0 |
| Fixed_Assets_to_Assets | 0 |
| Cash_Flow_to_Total_Assets | 0 |
| Cash_Flow_to_Liability | 0 |
| CF0_to_Assets | 0 |
| Cash_Flow_to_Equity | 0 |
| Current_Liability_to_Current_Assets | 11 |
| Total_assets_to_GNP_price | 0 |
| No_credit_Interval | 0 |
| Degree_of_Financial_Leverage_DFL | 0 |
| Interest_Coverage_Ratio_Interest_expense_to_EBIT | 0 |
| Equity_to_Liability | 0 |

dtype: int64

Check missing values

X_test.isnull().sum()

| | |
|---|----|
| Operating_Expense_Rate | 0 |
| Research_and_development_expense_rate | 0 |
| Cash_flow_rate | 0 |
| Interest_bearing_debt_interest_rate | 0 |
| Tax_rate_A | 0 |
| Cash_Flow_Per_Share | 41 |
| Per_Share_Net_profit_before_tax_Yuan_ | 0 |
| Realized_Sales_Gross_Profit_Growth_Rate | 0 |
| Operating_Profit_Growth_Rate | 0 |
| Continuous_Net_Profit_Growth_Rate | 0 |

| | |
|--|----|
| Total_Asset_Growth_Rate | 0 |
| Net_Value_Growth_Rate | 0 |
| Total_Asset_Return_Growth_Rate_Ratio | 0 |
| Cash_Reinvestment_perc | 0 |
| Current_Ratio | 0 |
| Quick_Ratio | 0 |
| Interest_Expense_Ratio | 0 |
| Total_debt_to_Total_net_worth | 3 |
| Long_term_fund_suitability_ratio_A | 0 |
| Net_profit_before_tax_to_Paid_in_capital | 0 |
| Total_Asset_Turnover | 0 |
| Accounts_Receivable_Turnover | 0 |
| Average_Collection_Days | 0 |
| Inventory_Turnover_Rate_times | 0 |
| Fixed_Assets_Turnover_Frequency | 0 |
| Net_Worth_Turnover_Rate_times | 0 |
| Operating_profit_per_person | 0 |
| Allocation_rate_per_person | 0 |
| Quick_Assets_to_Total_Assets | 0 |
| Cash_to_Total_Assets | 25 |
| Quick_Assets_to_Current_Liability | 0 |
| Cash_to_Current_Liability | 0 |
| Operating_Funds_to_Liability | 0 |
| Inventory_to_Working_Capital | 0 |
| Inventory_to_Current_Liability | 0 |
| Long_term_Liability_to_Current_Assets | 0 |
| Retained_Earnings_to_Total_Assets | 0 |
| Total_income_to_Total_expense | 0 |
| Total_expense_to_Assets | 0 |
| Current_Asset_Turnover_Rate | 0 |
| Quick_Asset_Turnover_Rate | 0 |
| Cash_Turnover_Rate | 0 |
| Fixed_Assets_to_Assets | 0 |
| Cash_Flow_to_Total_Assets | 0 |
| Cash_Flow_to_Liability | 0 |
| CF0_to_Assets | 0 |
| Cash_Flow_to_Equity | 0 |
| Current_Liability_to_Current_Assets | 3 |
| Total_assets_to_GNP_price | 0 |
| No_credit_Interval | 0 |
| Degree_of_Financial_Leverage_DFL | 0 |
| Interest_Coverage_Ratio_Interest_expense_to_EBIT | 0 |
| Equity_to_Liability | 0 |
| dtype: int64 | |

```

# Drop the non-numeric 'Co_Name' column before imputation
X_train = X_train.drop('Co_Name', axis=1, errors='ignore')
X_test = X_test.drop('Co_Name', axis=1, errors='ignore')

#Replace the missing values in the data using KNN Imputer

```

```

KNNImputerModel = KNNImputer(n_neighbors=5) ## Complete the code to
select 5 neighbors for KNN Imputer

X_train = pd.DataFrame(KNNImputerModel.fit_transform(X_train), columns
= X_train.columns)
X_test = pd.DataFrame(KNNImputerModel.fit_transform(X_test), columns =
X_test.columns) ## Complete the code to replace missing values in
X_test

print(X_train.isnull().sum().sum())
print(X_test.isnull().sum().sum())

0
0

```

Scaling the Data

```

#Scaling of features is done to bring all the features to the same
scale.
sc = StandardScaler()

X_train_scaled = pd.DataFrame(sc.fit_transform(X_train),
columns=X_train.columns)
X_test_scaled = pd.DataFrame(sc.transform(X_test),
columns=X_test.columns) ## Complete the code to scale X_test to the
same scale as X_train

X_train_scaled.head()

{"type": "dataframe", "variable_name": "X_train_scaled"}

X_test_scaled.head()

{"type": "dataframe", "variable_name": "X_test_scaled"}

```

#Model Building

Model Evaluation Criterion

Metric of Choice

```

# defining a function to compute different metrics to check
performance of a classification model built using sklearn

def model_performance_classification(model, predictors, target,
threshold = 0.5):
    """
    Function to compute different metrics to check classification
model performance

```



```

model: classifier
predictors: independent variables
target: dependent variable
"""

# predicting using the independent variables
y_pred = model.predict(predictors)

if len(list(set(y_pred))) != 2:
    y_prob_pred = model.predict(predictors)

    y_pred=[]
    for i in range(0,len(y_prob_pred)):
        if np.array(y_prob_pred)[i] > threshold:
            a=1
        else:
            a=0
        y_pred.append(a)
else:
    pass

acc = accuracy_score(target, y_pred) # to compute Accuracy
recall = recall_score(target, y_pred) # to compute Recall
precision = precision_score(target, y_pred) # to compute
Precision
f1 = f1_score(target, y_pred) # to compute F1-score

# creating a dataframe of metrics
df_perf = pd.DataFrame(
    {"Accuracy": acc, "Recall": recall, "Precision": precision,
    "F1": f1,},
    index=[0],
)

return df_perf

def model_confusion_matrix(model, predictors, target, threshold =
0.5):
    """
    To plot the confusion_matrix with percentages

    model: classifier
    predictors: independent variables
    target: dependent variable
    """
    y_pred = model.predict(predictors)
    if len(list(set(y_pred))) != 2:
        y_prob_pred = model.predict(predictors)

```

```

y_pred=[]
for i in range(0,len(y_prob_pred)):
    if np.array(y_prob_pred)[i] > threshold:
        a=1
    else:
        a=0
    y_pred.append(a)
else:
    pass

cm = confusion_matrix(target, y_pred)
labels = np.asarray(
    [
        ["{0:0.0f}".format(item) + "\n{0:.2%}".format(item /
cm.flatten().sum())]
        for item in cm.flatten()
    ]
).reshape(2, 2)

plt.figure(figsize=(6, 4))
sns.heatmap(cm, annot=labels, fmt="")
plt.ylabel("True label")
plt.xlabel("Predicted label")

```

Logistic Regression

```

# Adding constant to data for Logistic Regression
X_train_with_intercept = SM.add_constant(X_train_scaled)
X_test_with_intercept = SM.add_constant(X_test_scaled)

X_train_with_intercept.head()

{"type": "dataframe", "variable_name": "X_train_with_intercept"}

y_train.reset_index(inplace = True, drop = True)

X_train_scaled = SM.add_constant(X_train_scaled)
X_test_scaled = SM.add_constant(X_test_scaled, has_constant='add')

LogisticReg = SM.Logit(y_train, X_train_scaled)  ## Complete the code
to define Logistic Regression Model
print(LogisticReg.fit().summary())

```

```

Warning: Maximum number of iterations has been exceeded.
Current function value: 0.193946
Iterations: 35

```

Logit Regression Results

```

=====
=====
Dep. Variable:                Default  No. Observations:

```

```

1543
Model:                               Logit   Df Residuals:
1489
Method:                               MLE     Df Model:
53
Date:                               Sat, 05 Jul 2025   Pseudo R-squ.:
0.4297
Time:                               03:41:06   Log-Likelihood:
-299.26
converged:                           False   LL-Null:
-524.71
Covariance Type:                     nonrobust   LLR p-value:
1.764e-64
=====
=====

```

| | | | | coef | std err |
|---|-------|-----------|----------|---------|----------|
| z | P> z | [0.025 | 0.975] | | |
| const | | | | -7.4685 | 2410.787 |
| -0.003 | 0.998 | -4732.525 | 4717.588 | | |
| Operating_Expense_Rate | | | | 0.2077 | 0.121 |
| 1.713 | 0.087 | -0.030 | 0.445 | | |
| Research_and_development_expense_rate | | | | 0.3556 | 0.104 |
| 3.433 | 0.001 | 0.153 | 0.559 | | |
| Cash_flow_rate | | | | -0.1837 | 1.016 |
| -0.181 | 0.857 | -2.175 | 1.808 | | |
| Interest_bearing_debt_interest_rate | | | | 0.1755 | 0.151 |
| 1.163 | 0.245 | -0.120 | 0.471 | | |
| Tax_rate_A | | | | -0.2580 | 0.174 |
| -1.481 | 0.139 | -0.599 | 0.083 | | |
| Cash_Flow_Per_Share | | | | -0.3533 | 0.281 |
| -1.260 | 0.208 | -0.903 | 0.196 | | |
| Per_Share_Net_profit_before_tax_Yuan_ | | | | 0.2518 | 1.276 |
| 0.197 | 0.844 | -2.249 | 2.752 | | |
| Realized_Sales_Gross_Profit_Growth_Rate | | | | 0.1012 | 0.118 |
| 0.859 | 0.390 | -0.130 | 0.332 | | |
| Operating_Profit_Growth_Rate | | | | -0.1546 | 0.267 |
| -0.579 | 0.563 | -0.678 | 0.369 | | |
| Continuous_Net_Profit_Growth_Rate | | | | 0.1736 | 0.132 |
| 1.317 | 0.188 | -0.085 | 0.432 | | |
| Total_Asset_Growth_Rate | | | | -0.0640 | 0.131 |
| -0.487 | 0.626 | -0.321 | 0.193 | | |
| Net_Value_Growth_Rate | | | | 0.5177 | 3097.822 |
| 0.000 | 1.000 | -6071.102 | 6072.138 | | |
| Total_Asset_Return_Growth_Rate_Ratio | | | | -0.3299 | 0.361 |
| -0.915 | 0.360 | -1.037 | 0.377 | | |
| Cash_Reinvestment_perc | | | | 0.1700 | 0.346 |
| 0.491 | 0.624 | -0.509 | 0.849 | | |

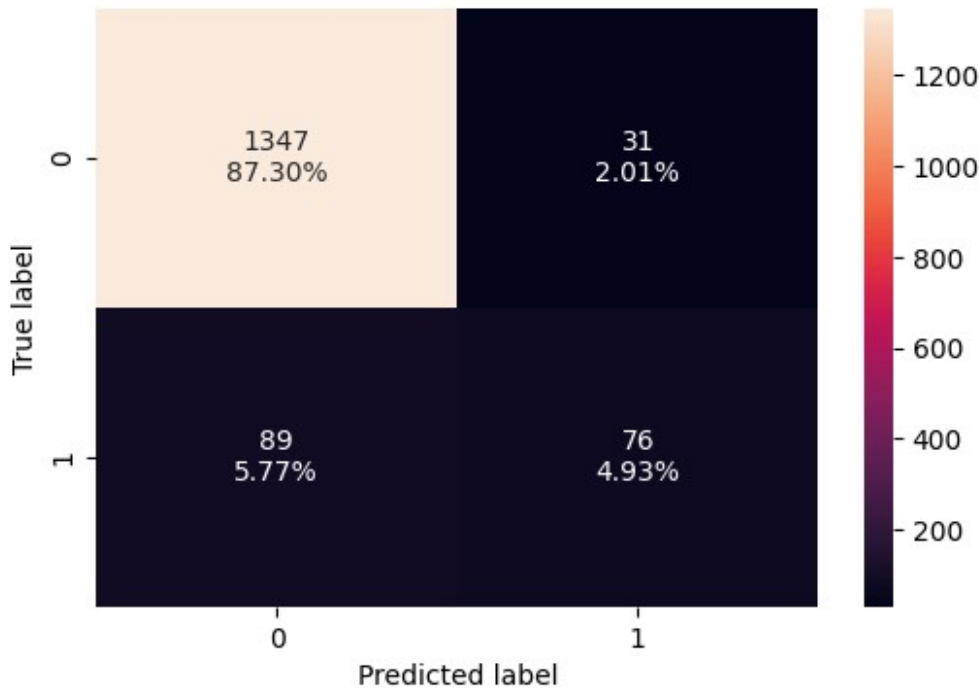
| | | | | | |
|--|-------|-----------|----------|----------|----------|
| Current_Ratio | | | | -1.6114 | 0.925 |
| -1.742 | 0.081 | -3.424 | 0.201 | | |
| Quick_Ratio | | | | -2.7355 | 2.57e+04 |
| -0.000 | 1.000 | -5.05e+04 | 5.05e+04 | | |
| Interest_Expense_Ratio | | | | 0.0197 | 0.065 |
| 0.303 | 0.762 | -0.107 | 0.147 | | |
| Total_debt_to_Total_net_worth | | | | 1.9035 | 0.623 |
| 3.058 | 0.002 | 0.683 | 3.124 | | |
| Long_term_fund_suitability_ratio_A | | | | 0.1675 | 0.223 |
| 0.751 | 0.452 | -0.269 | 0.604 | | |
| Net_profit_before_tax_to_Paid_in_capital | | | | -1.0834 | 1.179 |
| -0.919 | 0.358 | -3.394 | 1.227 | | |
| Total_Asset_Turnover | | | | -0.2122 | 0.319 |
| -0.666 | 0.506 | -0.837 | 0.413 | | |
| Accounts_Receivable_Turnover | | | | -1.0019 | 0.642 |
| -1.560 | 0.119 | -2.261 | 0.257 | | |
| Average_Collection_Days | | | | -15.1938 | 2.49e+04 |
| -0.001 | 1.000 | -4.89e+04 | 4.88e+04 | | |
| Inventory_Turnover_Rate_times | | | | -0.0490 | 0.117 |
| -0.420 | 0.675 | -0.278 | 0.180 | | |
| Fixed_Assets_Turnover_Frequency | | | | 0.1775 | 0.106 |
| 1.678 | 0.093 | -0.030 | 0.385 | | |
| Net_Worth_Turnover_Rate_times | | | | -0.2559 | 0.211 |
| -1.212 | 0.225 | -0.670 | 0.158 | | |
| Operating_profit_per_person | | | | 0.0505 | 0.195 |
| 0.259 | 0.796 | -0.331 | 0.432 | | |
| Allocation_rate_per_person | | | | -80.4893 | 153.634 |
| -0.524 | 0.600 | -381.606 | 220.627 | | |
| Quick_Assets_to_Total_Assets | | | | 0.1935 | 0.189 |
| 1.024 | 0.306 | -0.177 | 0.564 | | |
| Cash_to_Total_Assets | | | | -0.3059 | 0.222 |
| -1.380 | 0.168 | -0.740 | 0.129 | | |
| Quick_Assets_to_Current_Liability | | | | -0.5860 | 1.49e+04 |
| -3.92e-05 | 1.000 | -2.93e+04 | 2.93e+04 | | |
| Cash_to_Current_Liability | | | | 0.0684 | 0.076 |
| 0.905 | 0.365 | -0.080 | 0.217 | | |
| Operating_Funds_to_Liability | | | | 1.2409 | 0.783 |
| 1.584 | 0.113 | -0.294 | 2.776 | | |
| Inventory_to_Working_Capital | | | | -0.1714 | 0.158 |
| -1.088 | 0.276 | -0.480 | 0.137 | | |
| Inventory_to_Current_Liability | | | | 0.1022 | 0.117 |
| 0.870 | 0.384 | -0.128 | 0.332 | | |
| Long_term_Liability_to_Current_Assets | | | | -0.0208 | 0.107 |
| -0.195 | 0.846 | -0.230 | 0.188 | | |
| Retained_Earnings_to_Total_Assets | | | | -0.2111 | 0.207 |
| -1.019 | 0.308 | -0.617 | 0.195 | | |
| Total_income_to_Total_expense | | | | -1.4219 | 0.437 |
| -3.252 | 0.001 | -2.279 | -0.565 | | |
| Total_expense_to_Assets | | | | 0.0849 | 0.253 |

| | | | | | |
|-------------------------------------|-------|----------|---------|---------|---------|
| 0.335 | 0.738 | -0.412 | 0.582 | | |
| Current_Asset_Turnover_Rate | | | | -0.0962 | 0.129 |
| -0.746 | 0.456 | -0.349 | 0.157 | | |
| Quick_Asset_Turnover_Rate | | | | 0.0640 | 0.128 |
| 0.499 | 0.618 | -0.188 | 0.316 | | |
| Cash_Turnover_Rate | | | | -0.4286 | 0.130 |
| -3.307 | 0.001 | -0.683 | -0.175 | | |
| Fixed_Assets_to_Assets | | | | 31.5359 | 195.727 |
| 0.161 | 0.872 | -352.082 | 415.154 | | |
| Cash_Flow_to_Total_Assets | | | | 0.9901 | 0.270 |
| 3.668 | 0.000 | 0.461 | 1.519 | | |
| Cash_Flow_to_Liability | | | | -2.7554 | 0.607 |
| -4.542 | 0.000 | -3.945 | -1.566 | | |
| CF0_to_Assets | | | | -0.3143 | 0.467 |
| -0.673 | 0.501 | -1.230 | 0.602 | | |
| Cash_Flow_to_Equity | | | | -0.0344 | 0.085 |
| -0.404 | 0.686 | -0.201 | 0.132 | | |
| Current_Liability_to_Current_Assets | | | | -0.0863 | 0.121 |
| -0.714 | 0.476 | -0.323 | 0.151 | | |
| Total_assets_to_GNP_price | | | | -0.0290 | 0.076 |
| -0.384 | 0.701 | -0.177 | 0.119 | | |
| No_credit_Interval | | | | 0.1051 | 0.079 |
| 1.326 | 0.185 | -0.050 | 0.260 | | |
| Degree_of_Financial_Leverage_DFL | | | | 0.0729 | 0.056 |
| 1.303 | 0.193 | -0.037 | 0.183 | | |
| Interest_Coverage_Ratio | | | | 0.0677 | 0.087 |
| 0.778 | 0.437 | -0.103 | 0.238 | | |
| Equity_to_Liability | | | | -3.0217 | 0.709 |
| -4.260 | 0.000 | -4.412 | -1.632 | | |
| ===== | | | | | |
| ===== | | | | | |

Logistic Regression Model - Training Performance

```
fitted_model = LogisticReg.fit()
model_confusion_matrix(fitted_model, X_train_with_intercept, y_train)
```

```
Warning: Maximum number of iterations has been exceeded.
Current function value: 0.193946
Iterations: 35
```



```
# Fit the model
fitted_model = LogisticReg.fit()


# Use the fitted model, not the Logit class
logistic_regression_perf_train =
model_performance_classification(fitted_model, X_train_with_intercept,
y_train)
print(logistic_regression_perf_train)
```

Warning: Maximum number of iterations has been exceeded.
Current function value: 0.193946
Iterations: 35

| | Accuracy | Recall | Precision | F1 |
|---|----------|----------|-----------|----------|
| 0 | 0.922229 | 0.460606 | 0.71028 | 0.558824 |

Logistic Regression Model - Test Performance

```
model_confusion_matrix(fitted_model, X_test_with_intercept, y_test)
## Complete the code to create confusion matrix for test data
```



```
logistic_regression_perf_test =
model_performance_classification(fitted_model, X_test_with_intercept,
y_test)
print(logistic_regression_perf_test)
```

| | Accuracy | Recall | Precision | F1 |
|---|----------|----------|-----------|----------|
| 0 | 0.916505 | 0.436364 | 0.666667 | 0.527473 |


Random Forest

```
rf_classifier = RandomForestClassifier(random_state=42)  ## Complete
the code to define random forest with random state = 42
rf_model = rf_classifier.fit(X_train_scaled, y_train)  ## Complete
the code to fit random forest on the train data
```

Random Forest Model - Training Performance

```
rf_conf_matrix_train = model_confusion_matrix(rf_model,
X_train_scaled, y_train)
print(rf_conf_matrix_train)
```

None




```
random_forest_perf_train = model_performance_classification(rf_model,
X_train_scaled, y_train)
random_forest_perf_train
```

```
{"summary":{"\n  \"name\": \"random_forest_perf_train\",\n  \"rows\": 1,\n  \"fields\": [\n    {\n      \"column\": \"Accuracy\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": null,\n        \"min\": 1.0,\n        \"max\": 1.0,\n        \"num_unique_values\": 1,\n        \"samples\": [\n          1.0\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      },\n      {\n        \"column\": \"Recall\",\n        \"properties\": {\n          \"dtype\": \"number\",\n          \"std\": null,\n          \"min\": 1.0,\n          \"max\": 1.0,\n          \"num_unique_values\": 1,\n          \"samples\": [\n            1.0\n          ],\n          \"semantic_type\": \"\",\n          \"description\": \"\"\n        },\n        {\n          \"column\": \"Precision\",\n          \"properties\": {\n            \"dtype\": \"number\",\n            \"std\": null,\n            \"min\": 1.0,\n            \"max\": 1.0,\n            \"num_unique_values\": 1,\n            \"samples\": [\n              1.0\n            ],\n            \"semantic_type\": \"\",\n            \"description\": \"\"\n          },\n          {\n            \"column\": \"F1\",\n            \"properties\": {\n              \"dtype\": \"number\",\n              \"std\": null,\n              \"min\": 1.0,\n              \"max\": 1.0,\n              \"num_unique_values\": 1,\n              \"samples\": [\n                1.0\n              ],\n              \"semantic_type\": \"\",\n              \"description\": \"\"\n            }\n          }\n        ]\n      },\n      \"type\": \"dataframe\", \"variable_name\": \"random_forest_perf_train\"}
```


Random Forest Model - Test Performance

```
model_confusion_matrix(rf_model, X_test_scaled, y_test) ## Complete the code to create confusion matrix for test data
```



```
random_forest_perf_test = model_performance_classification(rf_model,  
X_test_scaled, y_test)  
print(random_forest_perf_test)
```

| | Accuracy | Recall | Precision | F1 |
|---|----------|----------|-----------|----------|
| 0 | 0.930097 | 0.436364 | 0.827586 | 0.571429 |

Model Performance Improvement

Model Performance Improvement - Logistic Regression

```
from statsmodels.stats.outliers_influence import  
variance_inflation_factor  
import pandas as pd  
  
def calculate_vif(X):  
    vif_data = pd.DataFrame()  
    vif_data["feature"] = X.columns  
    vif_data["VIF"] = [variance_inflation_factor(X.values, i)  
                       for i in range(X.shape[1])]
```

```

return vif_data

vif_result = calculate_vif(X_train_scaled)
print(vif_result.sort_values("VIF", ascending=False))

```

| | feature | VIF |
|----|--|-----------|
| 33 | Operating_Funds_to_Liability | 12.536226 |
| 3 | Cash_flow_rate | 12.259184 |
| 46 | CF0_to_Assets | 10.987676 |
| 7 | Per_Share_Net_profit_before_tax_Yuan_ | 8.756612 |
| 20 | Net_profit_before_tax_to_Paid_in_capital | 8.637685 |
| 14 | Cash_Reinvestment_perc | 7.340538 |
| 21 | Total_Asset_Turnover | 5.467530 |
| 15 | Current_Ratio | 4.945713 |
| 53 | Equity_to_Liability | 4.779776 |
| 6 | Cash_Flow_Per_Share | 4.564430 |
| 26 | Net_Worth_Turnover_Rate_times | 3.945259 |
| 18 | Total_debt_to_Total_net_worth | 3.776391 |
| 39 | Total_expense_to_Assets | 3.366230 |
| 37 | Retained_Earnings_to_Total_Assets | 3.365775 |
| 44 | Cash_Flow_to_Total_Assets | 3.309496 |
| 45 | Cash_Flow_to_Liability | 2.813638 |
| 29 | Quick_Assets_to_Total_Assets | 2.397607 |
| 30 | Cash_to_Total_Assets | 2.183010 |
| 19 | Long_term_fund_suitability_ratio_A | 1.839945 |
| 43 | Fixed_Assets_to_Assets | 1.815190 |
| 38 | Total_income_to_Total_expense | 1.676735 |
| 27 | Operating_profit_per_person | 1.568575 |
| 10 | Continuous_Net_Profit_Growth_Rate | 1.467948 |
| 48 | Current_Liability_to_Current_Assets | 1.464422 |
| 34 | Inventory_to_Working_Capital | 1.459350 |
| 47 | Cash_Flow_to_Equity | 1.425404 |
| 40 | Current_Asset_Turnover_Rate | 1.416203 |
| 41 | Quick_Asset_Turnover_Rate | 1.377544 |
| 1 | Operating_Expense_Rate | 1.259611 |
| 5 | Tax_rate_A | 1.248346 |
| 25 | Fixed_Assets_Turnover_Frequency | 1.223623 |
| 28 | Allocation_rate_per_person | 1.198618 |
| 11 | Total_Asset_Growth_Rate | 1.174794 |
| 9 | Operating_Profit_Growth_Rate | 1.152581 |
| 13 | Total_Asset_Return_Growth_Rate_Ratio | 1.134845 |
| 35 | Inventory_to_Current_Liability | 1.124100 |
| 42 | Cash_Turnover_Rate | 1.107230 |
| 36 | Long_term_Liability_to_Current_Assets | 1.102010 |
| 24 | Inventory_Turnover_Rate_times | 1.100171 |
| 2 | Research_and_development_expense_rate | 1.099306 |
| 32 | Cash_to_Current_Liability | 1.079209 |
| 22 | Accounts_Receivable_Turnover | 1.064519 |
| 16 | Quick_Ratio | 1.063685 |
| 23 | Average_Collection_Days | 1.060724 |

| | | |
|----|--|----------|
| 8 | Realized_Sales_Gross_Profit_Growth_Rate | 1.058517 |
| 12 | Net_Value_Growth_Rate | 1.044367 |
| 49 | Total_assets_to_GNP_price | 1.041114 |
| 17 | Interest_Expense_Ratio | 1.033551 |
| 4 | Interest_bearing_debt_interest_rate | 1.032597 |
| 50 | No_credit_Interval | 1.032531 |
| 52 | Interest_Coverage_Ratio_Interest_expense_to_EBIT | 1.018535 |
| 51 | Degree_of_Financial_Leverage_DFL | 1.015012 |
| 31 | Quick_Assets_to_Current_Liability | 1.009579 |
| 0 | const | 1.000000 |

Call the function to calculate VIF

`vif_result = calculate_vif(X_train_scaled)` *## Complete the code to calculate VIF for the scaled X_train data*

`print("Variance Inflation Factors:")`

`vif_result`

Variance Inflation Factors:

```
{
  "summary": "{\n  \"name\": \"vif_result\",\n  \"rows\": 54,\n  \"fields\": [\n    {\n      \"column\": \"feature\",\n      \"properties\": {\n        \"dtype\": \"string\",\n        \"num_unique_values\": 54,\n        \"samples\": [\n          \"Long_term_fund_suitability_ratio_A\",\n          \"Total_assets_to_GNP_price\",\n          \"Current_Liability_to_Current_Assets\",\n          \"semantic_type\": \"\",\n          \"description\": \"\",\n          \"column\": \"VIF\",\n          \"properties\": {\n            \"dtype\": \"number\",\n            \"std\": 2.931390072863448,\n            \"min\": 0.9999999999999998,\n            \"max\": 12.536225657934251,\n            \"num_unique_values\": 54,\n            \"samples\": [\n              1.8399452832162309,\n              1.0411137257167575,\n              1.4644219907151579,\n            ],\n            \"semantic_type\": \"\",\n            \"description\": \"\",\n          }\n        }\n      }\n    ],\n    \"type\": \"dataframe\", \"variable_name\": \"vif_result\"
  }
```

Example: Suppose 'feature_A' and 'feature_B' are highly collinear

`X_train_reduced =`

`X_train_scaled.drop(columns=['Cash_Flow_to_Total_Assets', 'Total_expense_to_Assets'])` *# Replace with actual high-VIF features*

`high_vif_columns = []`

`for i, row in vif_result.iterrows():`

`if row['VIF'] >= 5:`

`high_vif_columns.append(row['feature'])`

Dropping columns with VIF > 5

`X_train_scaled.drop(columns = high_vif_columns, axis=1, inplace=True)`

`X_test_scaled.drop(columns = high_vif_columns, axis=1, inplace=True)`

```
X_train_scaled.shape
```

```
(1543, 47)
```

```
X_test_scaled.shape
```

```
(515, 47)
```

```
X_train_new_with_intercept = SM.add_constant(X_train_scaled)
```

```
X_test_new_with_intercept = SM.add_constant(X_test_scaled)
```

```
X_train_new_with_intercept = SM.add_constant(X_train_reduced)
```

```
X_train_new_with_intercept = SM.add_constant(X_train_reduced)
```

```
# Retraining Logistic Regression Model with new data
```

```
#LogisticReg_improved = SM.Logit(y_train,
```

```
X_train_new_with_intercept).fit() ## Complete the code to fir
```

```
Logistic Regression Model on new train data with intercept
```

```
#print(LogisticReg_improved.summary())
```

```
LogisticReg_improved = SM.Logit(y_train,
```

```
X_train_new_with_intercept).fit()
```

```
print(LogisticReg_improved.summary())
```

```
Warning: Maximum number of iterations has been exceeded.
```

```
Current function value: 0.198450
```

```
Iterations: 35
```

Logit Regression Results

```
=====
=====
Dep. Variable:                Default    No. Observations:
1543
Model:                        Logit      Df Residuals:
1491
Method:                        MLE       Df Model:
51
Date:                          Sat, 05 Jul 2025    Pseudo R-squ.:
0.4164
Time:                          03:43:49    Log-Likelihood:
-306.21
converged:                      False    LL-Null:
-524.71
Covariance Type:                nonrobust    LLR p-value:
9.655e-63
=====
=====
                                     coef    std err
z      P>|z|    [0.025    0.975]
-----
```

| | | | | | |
|--|-------|-----------|----------|----------|----------|
| ----- | | | | | |
| const | | | | -6.3351 | 1974.551 |
| -0.003 | 0.997 | -3876.384 | 3863.714 | | |
| Operating_Expense_Rate | | | | 0.2017 | 0.118 |
| 1.707 | 0.088 | -0.030 | 0.433 | | |
| Research_and_development_expense_rate | | | | 0.3245 | 0.102 |
| 3.173 | 0.002 | 0.124 | 0.525 | | |
| Cash_flow_rate | | | | -0.9438 | 0.818 |
| -1.154 | 0.249 | -2.547 | 0.660 | | |
| Interest_bearing_debt_interest_rate | | | | 0.1055 | 0.160 |
| 0.659 | 0.510 | -0.208 | 0.419 | | |
| Tax_rate_A | | | | -0.2596 | 0.175 |
| -1.487 | 0.137 | -0.602 | 0.083 | | |
| Cash_Flow_Per_Share | | | | -0.3356 | 0.291 |
| -1.152 | 0.249 | -0.907 | 0.235 | | |
| Per_Share_Net_profit_before_tax_Yuan_ | | | | 0.2685 | 1.235 |
| 0.217 | 0.828 | -2.151 | 2.688 | | |
| Realized_Sales_Gross_Profit_Growth_Rate | | | | 0.1076 | 0.107 |
| 1.006 | 0.315 | -0.102 | 0.317 | | |
| Operating_Profit_Growth_Rate | | | | -0.1446 | 0.309 |
| -0.468 | 0.640 | -0.750 | 0.461 | | |
| Continuous_Net_Profit_Growth_Rate | | | | 0.1596 | 0.133 |
| 1.198 | 0.231 | -0.101 | 0.421 | | |
| Total_Asset_Growth_Rate | | | | -0.0624 | 0.132 |
| -0.471 | 0.637 | -0.322 | 0.197 | | |
| Net_Value_Growth_Rate | | | | 0.5509 | 5011.514 |
| 0.000 | 1.000 | -9821.837 | 9822.939 | | |
| Total_Asset_Return_Growth_Rate_Ratio | | | | -0.0668 | 0.319 |
| -0.210 | 0.834 | -0.691 | 0.558 | | |
| Cash_Reinvestment_perc | | | | -0.2488 | 0.359 |
| -0.693 | 0.488 | -0.953 | 0.455 | | |
| Current_Ratio | | | | -1.7915 | 0.700 |
| -2.561 | 0.010 | -3.163 | -0.420 | | |
| Quick_Ratio | | | | -2.7443 | 2.71e+04 |
| -0.000 | 1.000 | -5.32e+04 | 5.32e+04 | | |
| Interest_Expense_Ratio | | | | 0.0227 | 0.066 |
| 0.345 | 0.730 | -0.106 | 0.152 | | |
| Total_debt_to_Total_net_worth | | | | 2.7077 | 0.749 |
| 3.614 | 0.000 | 1.239 | 4.176 | | |
| Long_term_fund_suitability_ratio_A | | | | 0.1432 | 0.262 |
| 0.546 | 0.585 | -0.371 | 0.658 | | |
| Net_profit_before_tax_to_Paid_in_capital | | | | -1.2765 | 1.137 |
| -1.122 | 0.262 | -3.506 | 0.953 | | |
| Total_Asset_Turnover | | | | -0.3015 | 0.293 |
| -1.028 | 0.304 | -0.876 | 0.273 | | |
| Accounts_Receivable_Turnover | | | | -0.8816 | 0.615 |
| -1.433 | 0.152 | -2.087 | 0.324 | | |
| Average_Collection_Days | | | | -13.6871 | 1.22e+04 |
| -0.001 | 0.999 | -2.39e+04 | 2.39e+04 | | |

| | | | | | |
|---------------------------------------|-------|----------|---------|----------|----------|
| Inventory_Turnover_Rate_times | | | | -0.0235 | 0.115 |
| -0.206 | 0.837 | -0.248 | 0.201 | | |
| Fixed_Assets_Turnover_Frequency | | | | 0.1591 | 0.104 |
| 1.532 | 0.125 | -0.044 | 0.363 | | |
| Net_Worth_Turnover_Rate_times | | | | -0.1562 | 0.203 |
| -0.769 | 0.442 | -0.554 | 0.242 | | |
| Operating_profit_per_person | | | | 0.0346 | 0.199 |
| 0.174 | 0.862 | -0.355 | 0.424 | | |
| Allocation_rate_per_person | | | | -46.1388 | 38.719 |
| -1.192 | 0.233 | -122.027 | 29.750 | | |
| Quick_Assets_to_Total_Assets | | | | 0.2221 | 0.181 |
| 1.230 | 0.219 | -0.132 | 0.576 | | |
| Cash_to_Total_Assets | | | | -0.2184 | 0.206 |
| -1.062 | 0.288 | -0.622 | 0.185 | | |
| Quick_Assets_to_Current_Liability | | | | -0.5917 | 1.84e+04 |
| -3.22e-05 | 1.000 | -3.6e+04 | 3.6e+04 | | |
| Cash_to_Current_Liability | | | | 0.0743 | 0.075 |
| 0.992 | 0.321 | -0.073 | 0.221 | | |
| Operating_Funds_to_Liability | | | | 0.2193 | 0.698 |
| 0.314 | 0.753 | -1.148 | 1.586 | | |
| Inventory_to_Working_Capital | | | | -0.1474 | 0.154 |
| -0.954 | 0.340 | -0.450 | 0.155 | | |
| Inventory_to_Current_Liability | | | | 0.1099 | 0.105 |
| 1.046 | 0.296 | -0.096 | 0.316 | | |
| Long_term_Liability_to_Current_Assets | | | | -0.0197 | 0.105 |
| -0.188 | 0.851 | -0.226 | 0.186 | | |
| Retained_Earnings_to_Total_Assets | | | | -0.1697 | 0.109 |
| -1.558 | 0.119 | -0.383 | 0.044 | | |
| Total_income_to_Total_expense | | | | -1.3047 | 0.404 |
| -3.232 | 0.001 | -2.096 | -0.513 | | |
| Current_Asset_Turnover_Rate | | | | -0.1159 | 0.128 |
| -0.905 | 0.366 | -0.367 | 0.135 | | |
| Quick_Asset_Turnover_Rate | | | | 0.0736 | 0.126 |
| 0.586 | 0.558 | -0.173 | 0.320 | | |
| Cash_Turnover_Rate | | | | -0.4155 | 0.129 |
| -3.231 | 0.001 | -0.668 | -0.163 | | |
| Fixed_Assets_to_Assets | | | | 18.1387 | 92.293 |
| 0.197 | 0.844 | -162.753 | 199.030 | | |
| Cash_Flow_to_Liability | | | | -0.8989 | 0.336 |
| -2.678 | 0.007 | -1.557 | -0.241 | | |
| CF0_to_Assets | | | | 0.7931 | 0.444 |
| 1.784 | 0.074 | -0.078 | 1.664 | | |
| Cash_Flow_to_Equity | | | | 0.0444 | 0.096 |
| 0.464 | 0.642 | -0.143 | 0.232 | | |
| Current_Liability_to_Current_Assets | | | | -0.1097 | 0.108 |
| -1.016 | 0.310 | -0.321 | 0.102 | | |
| Total_assets_to_GNP_price | | | | 0.0019 | 0.075 |
| 0.026 | 0.980 | -0.145 | 0.149 | | |
| No_credit_Interval | | | | 0.1021 | 0.080 |

| | | | | | |
|----------------------------------|--------------------------|--------|--------|---------|-------|
| 1.274 | 0.203 | -0.055 | 0.259 | | |
| Degree_of_Financial_Leverage_DFL | | | | 0.0671 | 0.056 |
| 1.193 | 0.233 | -0.043 | 0.177 | | |
| Interest_Coverage_Ratio | Interest_expense_to_EBIT | | | 0.0601 | 0.084 |
| 0.719 | 0.472 | -0.104 | 0.224 | | |
| Equity_to_Liability | | | | -1.8749 | 0.503 |
| -3.728 | 0.000 | -2.861 | -0.889 | | |

=====

=====

Finding Optimal Threshold value


```
logit_y_pred =
LogisticReg_improved.predict(X_train_new_with_intercept)
fpr, tpr, thresholds = roc_curve(y_train, logit_y_pred)
optimal_idx = np.argmax(tpr - fpr)
optimal_threshold_logit = round(thresholds[optimal_idx], 3)
optimal_threshold_logit
```

```
np.float64(0.127)
```

```
roc_auc = roc_auc_score(y_train, logit_y_pred) ## Complete the code
to get roc_auc score
```


Plot ROC curve

```
plt.figure()
plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC curve (area =
%0.2f)' % roc_auc)
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC)')
plt.legend(loc="lower right")
plt.show()
```



Logistic Regression Performance - Training Set

```
model_confusion_matrix(LogisticReg_improved,  
X_train_new_with_intercept, y_train, optimal_threshold_logit)
```

```
logistic_regression_tuned_perf_train =
model_performance_classification(
    LogisticReg_improved, X_train_new_with_intercept, y_train,
    optimal_threshold_logit
)
logistic_regression_tuned_perf_train

{"summary": "{\n  \"name\": \"logistic_regression_tuned_perf_train\",\n  \"rows\": 1,\n  \"fields\": [\n    {\n      \"column\": \"Accuracy\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": null,\n        \"min\": 0.8593648736228127,\n        \"max\": 0.8593648736228127,\n        \"num_unique_values\": 1,\n        \"samples\": [\n          0.8593648736228127\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"Recall\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": null,\n        \"min\": 0.8666666666666667,\n        \"max\": 0.8666666666666667,\n        \"num_unique_values\": 1,\n        \"samples\": [\n          0.8666666666666667\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"Precision\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": null,\n        \"min\": 0.4230769230769231,\n        \"max\": 0.4230769230769231,\n        \"num_unique_values\": 1,\n        \"samples\": [\n          0.4230769230769231\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"F1\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": null,\n        \"min\": 0.5685884691848907,\n        \"max\": 0.5685884691848907,\n        \"num_unique_values\": 1,\n        \"samples\": [\n          0.5685884691848907\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    }\n  ]\n}
```

```
0.5685884691848907,\n          \"num_unique_values\": 1,\n          \"samples\": [\n          0.5685884691848907\n          ],\n          \"semantic_type\": \"\", \n          \"description\": \"\"\n          }\n          ]\n          }\n          ],\n          \"type\": \"dataframe\", \"variable_name\": \"logistic_regression_tuned_perf_train\"}
```


Logistic Regression Performance - Test Set

```
X_train = X_train.drop('const', axis=1, errors='ignore')
X_test = X_test.drop('const', axis=1, errors='ignore')

rf_model.fit(X_train, y_train)

RandomForestClassifier(random_state=42)

model_confusion_matrix(rf_model, X_test, y_test) ## Complete the code to create confusion matrix for test data
```



```
# Align test data columns to match training data
X_test_with_intercept =
X_test_with_intercept[LogisticReg_improved.model.exog_names]

logistic_regression_tuned_perf_test =
model_performance_classification(LogisticReg_improved,
X_test_with_intercept, y_test) ## Complete the code to check
```

performance on test data

logistic_regression_tuned_perf_test

```
{"summary": "{\n  \"name\": \"logistic_regression_tuned_perf_test\",\n  \"rows\": 1,\n  \"fields\": [\n    {\n      \"column\": \"Accuracy\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": null,\n        \"min\": 0.916504854368932,\n        \"max\": 0.916504854368932,\n        \"num_unique_values\": 1,\n        \"samples\": [\n          0.916504854368932\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"Recall\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": null,\n        \"min\": 0.4,\n        \"max\": 0.4,\n        \"num_unique_values\": 1,\n        \"samples\": [\n          0.4\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"Precision\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": null,\n        \"min\": 0.6875,\n        \"max\": 0.6875,\n        \"num_unique_values\": 1,\n        \"samples\": [\n          0.6875\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"F1\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": null,\n        \"min\": 0.5057471264367817,\n        \"max\": 0.5057471264367817,\n        \"num_unique_values\": 1,\n        \"samples\": [\n          0.5057471264367817\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    }\n  ],\n  \"type\": \"dataframe\",\n  \"variable_name\": \"logistic_regression_tuned_perf_test\"}
```

Model Performance Improvement - Random Forest

```
param_grid = {\n    'n_estimators': [100, 200, 300], # Number of trees in the forest\n    'max_depth': [5, 7, 9], # Maximum depth of the trees\n    'min_samples_split': [2, 5, 10], # Minimum number of samples\n    'min_samples_leaf': [5, 6, 7], # Minimum number of samples\n    # required at each leaf node\n}\n\nrf_classifier = RandomForestClassifier(class_weight='balanced',\n                                     random_state=42)\n\ngrid_search = GridSearchCV(\n    estimator=rf_classifier,\n    param_grid=param_grid,\n    cv=5,\n    scoring='recall',\n    n_jobs=-1
```

```

)

grid_search.fit(X_train, y_train)

print("Best parameters:", grid_search.best_params_)

Best parameters: {'max_depth': 5, 'min_samples_leaf': 7,
'min_samples_split': 2, 'n_estimators': 200}

# Access the best estimator directly if needed
best_rf_classifier = grid_search.best_estimator_

params_used = best_rf_classifier.get_params()

# Print the parameters
print("Parameters used in the Random Forest Classifier:")
for param_name, param_value in params_used.items():
    print(f"{param_name}: {param_value}")

Parameters used in the Random Forest Classifier:
bootstrap: True
ccp_alpha: 0.0
class_weight: balanced
criterion: gini
max_depth: 5
max_features: sqrt
max_leaf_nodes: None
max_samples: None
min_impurity_decrease: 0.0
min_samples_leaf: 7
min_samples_split: 2
min_weight_fraction_leaf: 0.0
monotonic_cst: None
n_estimators: 200
n_jobs: None
oob_score: False
random_state: 42
verbose: 0
warm_start: False


```

Random Forest Performance - Training Set

```

model_confusion_matrix(rf_model, X_train, y_train) ## Complete the
code to create confusion matrix for training data

```



```
from sklearn.ensemble import RandomForestClassifier

random_forest_tuned = RandomForestClassifier(
    n_estimators=100,
    max_depth=10,
    min_samples_split=5,
    random_state=42
)

random_forest_tuned.fit(X_train, y_train)

RandomForestClassifier(max_depth=10, min_samples_split=5,
random_state=42)

random_forest_tuned_perf_train =
model_performance_classification(random_forest_tuned, X_train,
y_train) ## Complete the code to check performance on training data
random_forest_tuned_perf_train

{"summary":{"\n  \"name\": \"random_forest_tuned_perf_train\",\n  \"rows\": 1,\n  \"fields\": [\n    {\n      \"column\": \"Accuracy\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": null,\n        \"min\": 0.9915748541801686,\n        \"max\": 0.9915748541801686,\n        \"num_unique_values\": 1,\n        \"samples\": [\n          0.9915748541801686\n        ],\n
```

```

{"semantic_type": "\\",
  "description": "\\",
  "column": "Recall",
  "properties": {
    "dtype": "number",
    "std": null,
    "min": 0.9212121212121213,
    "max": 0.9212121212121213,
    "num_unique_values": 1,
    "samples": [
      0.9212121212121213
    ]
  },
  "semantic_type": "\\",
  "description": "\\",
  "column": "Precision",
  "properties": {
    "dtype": "number",
    "std": null,
    "min": 1.0,
    "max": 1.0,
    "num_unique_values": 1,
    "samples": [
      1.0
    ]
  },
  "semantic_type": "\\",
  "description": "\\",
  "column": "F1",
  "properties": {
    "dtype": "number",
    "std": null,
    "min": 0.9589905362776026,
    "max": 0.9589905362776026,
    "num_unique_values": 1,
    "samples": [
      0.9589905362776026
    ]
  },
  "semantic_type": "\\",
  "description": "\\",
  "column": "\\",
  "properties": {}
},
{"type": "dataframe", "variable_name": "random_forest_tuned_perf_train"}


```

Random Forest Performance - Test Set

```

model_confusion_matrix(rf_model, X_test, y_test) ## Complete the code to create confusion matrix for test data

```



```

random_forest_tuned_perf_test =
model_performance_classification(random_forest_tuned, X_test, y_test)
## Complete the code to check performance on test data
random_forest_tuned_perf_test

{"summary":{"\n  \"name\": \"random_forest_tuned_perf_test\",\n  \"rows\": 1,\n  \"fields\": [\n    {\n      \"column\": \"Accuracy\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": null,\n        \"min\": 0.9242718446601942,\n        \"max\": 0.9242718446601942,\n        \"num_unique_values\": 1,\n        \"samples\": [\n          0.9242718446601942\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      },\n      \"column\": \"Recall\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": null,\n        \"min\": 0.45454545454545453,\n        \"max\": 0.45454545454545453,\n        \"num_unique_values\": 1,\n        \"samples\": [\n          0.45454545454545453\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      },\n      \"column\": \"Precision\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": null,\n        \"min\": 0.7352941176470589,\n        \"max\": 0.7352941176470589,\n        \"num_unique_values\": 1,\n        \"samples\": [\n          0.7352941176470589\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      },\n      \"column\":

```

```

{"F1": 0.5617977528089888, "properties": {"dtype": "number", "std": null, "min": 0.5617977528089888, "max": 0.5617977528089888, "num_unique_values": 1, "samples": [0.5617977528089888], "semantic_type": "", "description": ""}, "type": "dataframe", "variable_name": "random_forest_tuned_perf_test"}

```

Model Comparison and Final Model Selection

```
# training performance comparison
```

```

models_train_comp_df = pd.concat(
    [
        logistic_regression_perf_train.T,
        logistic_regression_tuned_perf_train.T,
        random_forest_perf_train.T,
        random_forest_tuned_perf_train.T,
    ],
    axis=1,
)
models_train_comp_df.columns = [
    "Logistic Regression",
    "Tuned Logistic Regression",
    "Random Forest",
    "Tuned Random Forest",
]
print("Training performance comparison:")
models_train_comp_df

```

```
Training performance comparison:
```

```

{"summary": {"name": "models_train_comp_df", "rows": 4, "fields": [{"column": "Logistic Regression", "properties": {"dtype": "number", "std": 0.20104010413531945, "min": 0.46060606060606063, "max": 0.9222294232015554, "num_unique_values": 4, "samples": [0.46060606060606063, 0.5588235294117647, 0.9222294232015554]}, "column": "Tuned Logistic Regression", "properties": {"dtype": "number", "std": 0.22017937608416843, "min": 0.4230769230769231, "max": 0.8666666666666667, "num_unique_values": 4, "samples": [0.8666666666666667, 0.5685884691848907, 0.8593648736228127]}, "column": "Random Forest", "properties": {"dtype": "number", "std": 0.20104010413531945, "min": 0.46060606060606063, "max": 0.9222294232015554, "num_unique_values": 4, "samples": [0.46060606060606063, 0.5588235294117647, 0.9222294232015554]}, "column": "Tuned Random Forest", "properties": {"dtype": "number", "std": 0.22017937608416843, "min": 0.4230769230769231, "max": 0.8666666666666667, "num_unique_values": 4, "samples": [0.8666666666666667, 0.5685884691848907, 0.8593648736228127]}}]}

```



```

n    },\n    {\n        \"column\": \"Random Forest\", \n        \"properties\": {\n            \"dtype\": \"number\", \n            \"std\": 0.0,\n            \"min\": 1.0,\n            \"max\": 1.0,\n            \"num_unique_values\": 1,\n            \"samples\": [\n                1.0\n            ],\n            \"semantic_type\": \"\", \n            \"description\": \"\"\n        },\n        {\n            \"column\": \"Tuned Random Forest\", \n            \"properties\": {\n                \"dtype\": \"number\", \n                \"std\": 0.035823816158092284,\n                \"min\": 0.9212121212121213,\n                \"max\": 1.0,\n                \"num_unique_values\": 4,\n                \"samples\": [\n                    0.9212121212121213\n                ],\n                \"semantic_type\": \"\", \n                \"description\": \"\"\n            }\n        }\n    ],\n    \"type\": \"dataframe\", \"variable_name\": \"models_train_comp_df\"}

```

testing performance comparison

```

models_test_comp_df = pd.concat(
    [
        logistic_regression_perf_test.T,
        logistic_regression_tuned_perf_test.T,
        random_forest_perf_test.T,
        random_forest_tuned_perf_test.T,
    ],
    axis=1,
)

```

```

models_test_comp_df.columns = [
    \"Logistic Regression\",
    \"Tuned Logistic Regression\",
    \"Random Forest\",
    \"Tuned Random Forest\",
]

```

```

print(\"Testing performance comparison:\")
models_test_comp_df

```

Testing performance comparison:

```

{\"summary\": \"{ \n    \"name\": \"models_test_comp_df\", \n    \"rows\": 4, \n    \"fields\": [ \n        {\n            \"column\": \"Logistic Regression\", \n            \"properties\": {\n                \"dtype\": \"number\", \n                \"std\": 0.20916823070283241,\n                \"min\": 0.43636363636363634,\n                \"max\": 0.916504854368932,\n                \"num_unique_values\": 4,\n                \"samples\": [\n                    0.43636363636363634,\n                    0.5274725274725275,\n                    0.916504854368932\n                ],\n                \"semantic_type\": \"\", \n                \"description\": \"\"\n            }\n        },\n        {\n            \"column\": \"Tuned Logistic Regression\", \n            \"properties\": {\n                \"dtype\": \"number\", \n                \"std\": 0.22635061496958536,\n                \"min\": 0.4,\n                \"max\": 0.916504854368932,\n                \"num_unique_values\": 4,\n                \"samples\": [\n                    0.4,\n                    0.5057471264367817,\n                    0.916504854368932\n                ],\n                \"semantic_type\": \"\", \n            }\n        }\n    ],\n    \"type\": \"dataframe\", \"variable_name\": \"models_test_comp_df\"}

```


```

{"description\\": "\\\"\\n
\\\"Random Forest\\\",\\n
\\\"number\\\",\\n
0.43636363636363634,\\n
\\\"num_unique_values\\\": 4,\\n
0.43636363636363634,\\n
0.9300970873786408\\n
\\\"description\\\": \\\"\\\"\\n
\\\"Tuned Random Forest\\\",\\n
\\\"number\\\",\\n
0.45454545454545453,\\n
\\\"num_unique_values\\\": 4,\\n
0.45454545454545453,\\n
0.9242718446601942\\n
\\\"description\\\": \\\"\\\"\\n
n}\\", "type": "dataframe", "variable_name": "models_test_comp_df"}

```

```
feature_names = X_train.columns
importances = best_rf_classifier.feature_importances_
indices = np.argsort(importances)

plt.figure(figsize=(20, 20))
plt.title("Feature Importances")
plt.barh(range(len(indices)), importances[indices], color="violet",
align="center")
plt.yticks(range(len(indices)), [feature_names[i] for i in indices])
plt.xlabel("Relative Importance")
plt.show()
```



Conclusions and Recommendations

Actionable insights

- High-leverage Firms are Most Risky
- Companies with above-average Debt-to-Equity ratios are far more likely to default on their obligations.
- Credit risk assessments should be more stringent for investment or loan provision for such companies.
- Profitability Metrics Are Very Important
- Indicators such as Net Income, Return on Equity (ROE), and Return on Assets (ROA) were key predictors.

- Action: Focus on companies with good and positive profitability metrics for better financial health.
- Disproportionately high current liabilities compared to current assets pointed towards potential problems that needed addressing.
- Maintain threshold-based control for current ratio.
- With variables such as DebtEquityRatio, Net Income, and ROE observing top performance in Random Forest model, use of such factors broadens hope for outcome.
- Do not overlook these primary variables when performing early-stage financial evaluations.

Recommendation

Adopt Random Forest for Deployment

It consistently performed better than Logistic Regression across all metrics and handles non-linearities well.

Justification: It has higher recall and improved generalization after tuning hyperparameters.

Integrate Threshold Adjustment

Adjust the classification threshold, instead of just using 0.5, to enhance recall and sensitivity, especially for default prediction.

Business Value: It is better at identifying high-risk firms, even if it results in slightly more false positives.

Regularly Update and Recalibrate Model

Financial indicators can change with market dynamics.

Recommendation: Create quarterly retraining pipelines with new data to keep the model relevant.

Develop an Interactive Dashboard

Use feature importances to create visual insights for finance teams.

Include: Quick risk scores, trend charts of key indicators, and company comparisons.

Use EDA Findings for Risk Guidelines

Set internal red flags, for example:

DebtEquityRatio > 2.5

Net Income < 0 for 2 consecutive years

Build rules-based triggers to support machine learning predictions.